

# lab1

October 19, 2023

## 1 LAB 1 - Wykrywanie krawędzi metodą Canny’ego

Autor: Błażej Nowicki

```
[1]: import torch
import torch.nn as nn
import numpy as np
from PIL import Image, ImageOps
import torchvision.transforms as T
```

```
[2]: def preview(tensor):
    preview = tensor.squeeze(0)
    return T.ToPILImage()(preview)

def save(pil, name):
    return pil.save(f'images/{name}.png', "PNG")
```

### 1.0.1 Load image

```
[3]: img = original = Image.open("sample.jpg")
img
save(img, "0_original")
```

```
[4]: img = T.ToTensor()(img)
img = img.unsqueeze(0)  # == reshape(1, 2, w, h)
```

### 1.0.2 Convert RGB to grayscale

```
[5]: # 0.299 Red + 0.587 Green + 0.114 Blue.
def rgb_to_grayscale_kernel():
    kernel = torch.tensor([0.299, 0.587, 0.114], dtype=torch.float32)
    return kernel.reshape((1, 3, 1, 1))
```

```
[6]: kernel = rgb_to_grayscale_kernel()
rgb_to_gray_conv = nn.Conv2d(in_channels=3, out_channels=1, kernel_size=1,
    ↪ bias=False)
rgb_to_gray_conv.weight = nn.Parameter(kernel)
```

```
[7]: img = rgb_to_gray_conv(img)
```

```
[8]: pil = preview(img)
      save(pil, "1_gray_scale")
      pil
```

```
[8]:
```



### 1.0.3 Apply average pooling

Do wykrywania krawędzi zastosujemy MaxPool ponieważ nie rozmazuje krawędzi tak jak AvgPool. Przy pullingu gdzie kernel size = 2 różnica pomiędzy tymi dwoma jest prawie niezauważalna.

```
[9]: avg_poll_conv = nn.MaxPool2d(kernel_size=2)
      print(img.shape)
      img = avg_poll_conv(img)

      pil = preview(img)
      save(pil, "2_resized")
      pil
```

```
torch.Size([1, 1, 360, 540])
```

```
[9]:
```



#### 1.0.4 Apply gaussian blur

```
[10]: def gaussian_kernel(size, sigma=1):  
    assert size % 2 == 1, "Kernel size must be an odd number"  
    half = int(size) // 2  
    x, y = np.mgrid[-half : half + 1, -half : half + 1]  
    normal = 1 / (2.0 * np.pi * sigma**2)  
    g = np.exp(-((x**2 + y**2) / (2.0 * sigma**2))) * normal  
    return torch.tensor(g, dtype=torch.float32).reshape(1, 1, size, size)
```

```
[11]: blur_size = 21  
kernel = gaussian_kernel(blur_size)  
blur_conv = nn.Conv2d(  
    in_channels=1, out_channels=1, kernel_size=kernel.shape, bias=False  
)  
blur_conv.weight = nn.Parameter(kernel)  
img = blur_conv(img)
```

```
[12]: pil = preview(img)  
save(pil, "3_blured")  
pil
```

```
[12]:
```



### 1.0.5 Calculate gradient

```
[13]: def gradient_kernel():
        kernel = torch.tensor(
            [
                [[[[1, 2, 1], [0, 0, 0], [-1, -2, -1]]],
                 [[[[1, 0, -1], [2, 0, -2], [1, 0, -1]]],
            ],
            dtype=torch.float32,
        )
        return kernel.reshape(2, 1, 3, 3)

[14]: kernel = gradient_kernel()
print(kernel)
gradient_conv = nn.Conv2d(
    in_channels=1, out_channels=2, kernel_size=3, bias=False, #stride=(1, 1, 1, 1), padding=(0, 0, 0), dilation=(0, 0, 0)
)
gradient_conv.weight = nn.Parameter(kernel)
img = gradient_conv(img)


tensor([[[[ 1.,  2.,  1.],
           [ 0.,  0.,  0.],
           [-1., -2., -1.]],

         [[ 1.,  0., -1.],
           [ 2.,  0., -2.],
           [ 1.,  0., -1.]]]])

[15]: pil = preview(img[:, 0, :, :])
save(pil, "4_sobel_x")
```


```
pil
```

```
[15]:
```



```
[16]: pil = preview(img[:, 1, :, :])  
save(pil, "5_sobel_y")  
pil
```

```
[16]:
```



```
[17]: class Gradient(nn.Module):  
    def forward(self, x):  
        return torch.concat(  
            ((torch.sqrt(torch.sum(torch.square(x), dim=1))).unsqueeze(1),  
             torch.arctan(x[:, 0, :, :] / x[:, 1, :, :]).unsqueeze(1)),  
            dim=1,  
        )
```

```
[18]: gradient_layer = Gradient()  
img = gradient_layer(img)  
img.shape
```

```
[18]: torch.Size([1, 2, 158, 248])
```


```
[19]: pil = preview(img[:, 0, :, :])  
      save(pil, "6_gradient_value")  
      pil
```

```
[19]:
```



```
[20]: pil = preview(img[:, 1, :, :])  
      save(pil, "7_gradient_direction")  
      pil
```

```
[20]:
```



```
[21]: test = img  
      test.shape
```

```
[21]: torch.Size([1, 2, 158, 248])
```

```
[22]: class NonMaxSuppression(nn.Module):  
      def forward(self, img):  
          M, N = img.shape[-2:]
```

```

D = img[0, 1, :, :]
img = img[0, 0, :, :]
# return img

Z = np.zeros((M, N), dtype=np.float32)
angle = D * 180.0 / np.pi
angle[angle < 0] += 180
for i in range(1, M - 1):
    for j in range(1, N - 1):
        try:
            q = float("inf")
            r = float("inf")
            # angle 0
            if (0 <= angle[i, j] < 22.5) or (157.5 <= angle[i, j] <= 180):
                q = img[i, j + 1]
                r = img[i, j - 1]
            # angle 45
            elif 22.5 <= angle[i, j] < 67.5:
                q = img[i + 1, j - 1]
                r = img[i - 1, j + 1]
            # angle 90
            elif 67.5 <= angle[i, j] < 112.5:
                q = img[i + 1, j]
                r = img[i - 1, j]
            # angle 135
            elif 112.5 <= angle[i, j] < 157.5:
                q = img[i - 1, j - 1]
                r = img[i + 1, j + 1]
            if (img[i, j] >= q) and (img[i, j] >= r):
                Z[i, j] = img[i, j]
            else:
                Z[i, j] = 0
        except IndexError as e:
            pass

    return torch.tensor(Z, dtype=torch.float32).reshape(
        1, 1, Z.shape[-2], Z.shape[-1]
    )

```

```

[23]: non_max_suppression_layer = NonMaxSuppression()
img = non_max_suppression_layer(img)

pil = preview(img)
save(pil, "8_non_max_suppression")


```

```
[24]: TRESHHOLD = 0.5

class ReLU(nn.Module):
    def forward(self, x):
        x[x < TRESHHOLD] = 0.
        return x

relu = ReLU()
img = relu(img)
final = preview(img)
save(final, "9_relu")
final
```

[24]:



```
[25]: expanded = ImageOps.expand(final, border=12, fill="black")
resized = expanded.resize(original.size)
save(resized, "10_final_resized")
resized
```

[25]:





Żeby krawędzie pokrywały się z oryginalnym obrazkiem należy dodać marginesy na końcowy obrazek. Alternatywnie na każdym etapie obróbki trzeba by dbać o poprawny padding, żeby filtry konwolucyjne nie zmieniały wielkości obrazu.

```
[26]: overlaid = original.copy()
      overlaid.paste(resized, (0,0), resized)
      save(overlaid, "11_overlaid")
      overlaid
```

[26]:

