

gan

November 16, 2023

1 Lab 3 - GANy

Błażej Nowicki

```
[ ]: import torch
      from torch import nn

      import math
      import matplotlib.pyplot as plt
      import torchvision
      import torchvision.transforms as transforms
      from PIL import Image
      from torch.utils.data import DataLoader, Dataset
      import os

      torch.manual_seed(111)

      assert torch.cuda.is_available(), "Ten mały manewr kosztował by nas 10 lat"
      device = "cuda"

      checkpoint_dir = "checkpoints/"
```

1.0.1 Rozgrzewka

```
[ ]: class MiniModel(nn.Module):
      def __init__(self):
          super().__init__()
          self.model = nn.Sequential(
              nn.Linear(2, 10),
              nn.Linear(10, 1),
          )

      def forward(self, x):
          return self.model(x)

      batch = 3
```

```

model = MiniModel()
loss_function = lambda x: torch.abs(torch.tensor(42.0) - x).mean()
optimizer = torch.optim.Adam(list(model.parameters())[:2], lr=0.1) # 2 bo wagi
    ↪ + bias
epochs = 1000

x = torch.randn((batch, 2))

for e in range(epochs):
    model.zero_grad()
    y = model(x)
    loss = loss_function(y)
    loss.backward()
    optimizer.step()

    if (e + 1) % 100 == 0:
        print(loss)

model.eval()
model(x)

```

```

tensor(15.6990, grad_fn=<MeanBackward0>)
tensor(6.1553, grad_fn=<MeanBackward0>)
tensor(4.9915, grad_fn=<MeanBackward0>)
tensor(3.7765, grad_fn=<MeanBackward0>)
tensor(2.5117, grad_fn=<MeanBackward0>)
tensor(1.2427, grad_fn=<MeanBackward0>)
tensor(0.0512, grad_fn=<MeanBackward0>)
tensor(0.0449, grad_fn=<MeanBackward0>)
tensor(0.0060, grad_fn=<MeanBackward0>)
tensor(0.0462, grad_fn=<MeanBackward0>)

```

```

[ ]: tensor([[42.0558],
            [41.9856],
            [42.0294]], grad_fn=<AddmmBackward0>)

```

Model zbiega do oczekiwanej wartości >42<

```

[ ]: model.state_dict()

```

```

[ ]: OrderedDict([('model.0.weight',
                    tensor([[ -2.1645, -14.9439],
                          [ -2.4499, -16.0248],
                          [ -3.1218, -15.0403],
                          [  1.8771,  15.7340],
                          [  2.6101,  15.5759],
                          [  2.7396,  14.9644],
                          [ -2.1010, -14.7368],

```

```

        [ -1.8904, -15.8519],
        [  2.0407,  15.9046],
        [ -3.1293, -15.9088]])),
    ('model.0.bias',
     tensor([ 25.5340,  24.4872,  25.6582, -24.8000, -25.2096,
-24.5091,  24.5038,
           25.7281, -24.6011,  25.1103])),
    ('model.1.weight',
     tensor([[ 0.0988,  0.0381,  0.0973, -0.1788, -0.0612, -0.2741,
 0.1142,  0.0515,
           -0.0069,  0.2939]])),
    ('model.1.bias', tensor([0.2898])))

```

Zmieniły się tylko wagi i bias w pierwszej warstwie. W drugiej zostały wartości z inicjalizacji.

1.0.2 Modele

```

[ ]: class Discriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=4, stride=2,
↪padding=1),
            nn.BatchNorm2d(num_features=32),
            nn.LeakyReLU(0.2),
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=4, stride=2,
↪padding=1),
            nn.BatchNorm2d(num_features=64),
            nn.LeakyReLU(0.2),
            nn.Conv2d(in_channels=64, out_channels=64, kernel_size=4, stride=2,
↪padding=1),
            nn.BatchNorm2d(num_features=64),
            nn.LeakyReLU(0.2),
            nn.Flatten(),
            nn.Dropout(0.2),
            nn.Linear(1024, 1),
        )

    def forward(self, x):
        return self.model(x)

class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Linear(64, 1024),

```

```

nn.Unflatten(-1, (64, 4, 4)),
# ----
nn.Upsample(scale_factor=2),
nn.Conv2d(64, 64, 3, padding=1),
nn.BatchNorm2d(64),
nn.LeakyReLU(0.2),
nn.Upsample(scale_factor=2),
nn.Conv2d(64, 128, 3, padding=1),
nn.BatchNorm2d(128),
nn.LeakyReLU(0.2),
nn.Upsample(scale_factor=2),
nn.Conv2d(128, 256, 3, padding=1),
nn.BatchNorm2d(256),
nn.LeakyReLU(0.2),
nn.Conv2d(256, 256, 3, padding=1),
nn.BatchNorm2d(256),
nn.LeakyReLU(0.2),
# ---
# nn.ConvTranspose2d(64, 64, 4, stride=2, padding=1),
# nn.LeakyReLU(0.2),
# nn.BatchNorm2d(64),
# nn.ConvTranspose2d(64, 128, 4, stride=2, padding=1),
# nn.LeakyReLU(0.2),
# nn.BatchNorm2d(128),
# nn.ConvTranspose2d(128, 256, 4, stride=2, padding=1),
# nn.LeakyReLU(0.2),
# nn.BatchNorm2d(256),
# ---
nn.Conv2d(256, 3, 5, stride=1, padding=2),
nn.Tanh()

)

def forward(self, x):
    return self.model(x)

```

```


[ ]: def show(x):
      x = (x + 1) / 2
      plt.imshow(x.permute(1, 2, 0).cpu().detach().numpy())

```

```

[ ]: model = Generator()
      x = torch.randn(2, 64)
      x = model(x)
      show(x[0])

```



Bez zaskoczeń kolorowe plamy. Z `convTransposed2d` nie było płynnych przejść pomiędzy sąsiednimi pikselami.

```
[ ]: class CakeDataset(Dataset):
    def __init__(self, data_dir: str):
        transform = transforms.Compose([
            transforms.ToTensor(),
        ])
        self.images = []

        for file in os.listdir(data_dir):
            img = Image.open(data_dir + file)
            img = img.resize((32, 32))
            img = transform(img) # [0, 1]
            img = 2 * img - 1 # [-1, 1]
            self.images.append(img)

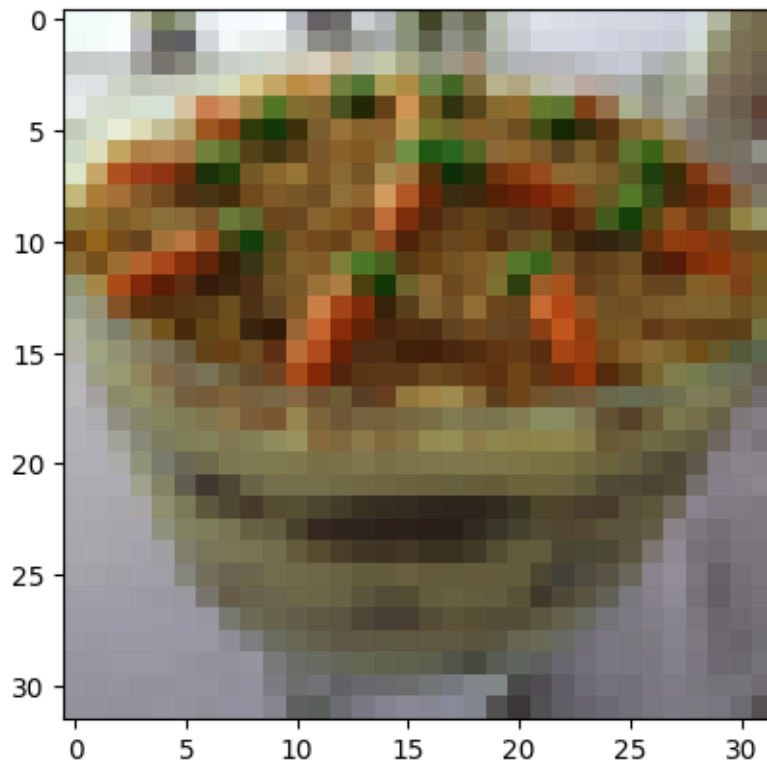
    def __len__(self):
        return len(self.images)

    def __getitem__(self, index) -> Image:
```

```
return self.images[index]
```

```
[ ]: dataset = CakeDataset('crawled_cakes/')
```

```
[ ]: show(dataset[1])
```



```
[ ]: batch_size = 64
train_loader = torch.utils.data.DataLoader(
    dataset, batch_size=batch_size, shuffle=True, drop_last=True
)

start_epoch = 0
end_epoch = 3000

loss_function = nn.MSELoss().to(device)
discriminator = Discriminator().to(device=device)
generator = Generator().to(device=device)

lr = 0.0001
optimizer_discriminator = torch.optim.Adam(discriminator.parameters(), lr=lr)
optimizer_generator = torch.optim.Adam(generator.parameters(), lr=lr)
```

```

if start_epoch > 0:
    checkpoint = torch.load(checkpoint_dir + f"model_{start_epoch-1}.pt")

    discriminator.load_state_dict(checkpoint['discriminator_weights'])
    generator.load_state_dict(checkpoint["generator_weights"])
    optimizer_discriminator.
    ↪load_state_dict(checkpoint['discriminator_optimizer'])
    optimizer_generator.load_state_dict(checkpoint['generator_optimizer'])

for epoch in range(start_epoch, end_epoch):
    for n, real_samples in enumerate(train_loader):
        # Data for training the discriminator
        real_samples = real_samples.to(device=device)
        real_samples_labels = torch.ones((batch_size, 1)).to(
            device=device
        )

        input_noise = torch.randn((batch_size, 64)).to(
            device=device
        )
        generated_samples = generator(input_noise)
        generated_samples_labels = torch.zeros((batch_size, 1)).
        ↪to(device=device)

        generator.zero_grad()
        output_discriminator_generated = discriminator(generated_samples)

        loss_generator = loss_function(
            output_discriminator_generated, real_samples_labels
        )
        loss_generator.backward()
        optimizer_generator.step()

        discriminator.zero_grad()
        output_discriminator_real = discriminator(real_samples)
        output_discriminator_generated = discriminator(generated_samples.
        ↪detach())

        loss_discriminator = loss_function(
            torch.cat([output_discriminator_real,
            ↪output_discriminator_generated]),
            torch.cat([real_samples_labels, generated_samples_labels])
        )

        if loss_discriminator > 0.2:
            loss_discriminator.backward()
            optimizer_discriminator.step()

```

```

# Show loss
if n == len(train_loader) - 1 :

    print(f"Epoch: {epoch} Loss D.: {loss_discriminator}")
    print(f"Epoch: {epoch} Loss G.: {loss_generator}")

if (epoch+1) % 100 == 0:
    path = checkpoint_dir + f"model_{epoch}.pt"
    torch.save({
        "generator_weights": generator.state_dict(),
        "generator_optimizer": optimizer_generator.state_dict(),
        "discriminator_weights": discriminator.state_dict(),
        "discriminator_optimizer": optimizer_discriminator.state_dict(),
        "epoch": epoch
    }, path)

```

```

[ ]: def show_epoch(x, epoch):
    checkpoint = torch.load(checkpoint_dir + f"model_{epoch-1}.pt")
    generator = Generator()
    generator.load_state_dict(checkpoint["generator_weights"])
    y = generator(x)

    y = (y + 1) / 2
    fig = plt.figure(figsize=(40, 20))
    columns = y.shape[0]
    for i in range(1, columns + 1):
        fig.add_subplot(1, columns, i)
        plt.imshow(y[i-1].permute(1, 2, 0).cpu().detach().numpy())
    plt.show()

```

```


[ ]: x = torch.randn((10, 64))

```

```

[ ]: show_epoch(x, 100)


```




```

[ ]: show_epoch(x, 200)


```





```
[ ]: show_epoch(x, 500)
```




```
[ ]: show_epoch(x, 1000)
```




```
[ ]: show_epoch(x, 1500)
```




```
[ ]: show_epoch(x, 2000)
```



```
[ ]: show_epoch(x, 2500)
```



```
[ ]: show_epoch(x, 3000)
```



Jak widać modele jak i parametry mocno oddryfowały od tych podanych w labie. Przez długi czas z domyślnymi nie mogłem uzyskać efektów inne niż różowe plamy. Z tymi modyfikacjami trening ruszył ale dalej były problemy. Na początku testowałem z ConvTransposed2d ze słabymi efektami dlatego zmieniłem na upscale + conv.

Widać że generator wykorzystuje błędy dyskriminaora i uczy się generować tylko kilka typów obrazków. Główny podejrzany to zwiększony learning rate